

Solving N-body Problem via Barnes- Hut Algorithm

GPU ACCELERATED COMPUTING


SYED MUHAMMAD
SUFFWAN

SYED MUHAMMAD HASAN
HAIDER

01


Abstract

- N-body problem is predicting the behaviour of a mass in presence of the other masses and the masses could range.
- The current computational cost of simulating of all-pairs interactions is $O(n^2)$.
- We parallelised Barnes-hut algorithm for N-body problem using CUDA
- Computational cost brought down to $O(n \log n)$.
- All these approaches have been implemented out on NVIDIA Tesla K80



Application - Barnes-Hut Algorithm


- **Approximation algorithm for n-body simulation**
- **Approximate a particle's dependence on distant data**
- **We use approximations on their centre of mass and total mass with mass of each body is assumed to be constant**
- **It uses a oct-tree which is similar to a binary tree, except that each node has 8 children (some of which may be empty).**



Application - Barnes-Hut Algorithm

- **Recursively divides the n bodies into groups while storing in an octree (3D simulation)**
- **Each node in the tree is 3D space.**
- **Breakdown the space until each division has a 1 or zero bodies.**
- **Internal Node: Represents group of bodies beneath it. Stores centre of mass of all its children**
- **External Node: No children (at consist one body at max)**
- **We have implemented this algorithm in three different ways; Serial, Naive and Shared Memory.**

an octree decomposition



Serial Implementation Bottlenecks




BUILDING OCT TREE



**TREE TRAVERSAL IN
FORCE COMPUTATION**


Serial Implementation Performance



Naive Implementation Performance



UPDATING POSITIONS



BARNES-HUT SHARED MEMORY EXPLANATION


EACH PROCESSOR FIRST BUILDS A LOCAL TREE; THESE ARE MERGED INTO A GLOBAL TREE STORED IN SHARED MEMORY. WORK IS EVENLY DISTRIBUTED AMONG PROCESSORS BY PARTITIONING THE BODIES.

EACH BLOCK HAS EQUAL CHUNK OF DATA

EACH BLOCK FINDS THE LOCAL MINIMA AND MAXIMA.

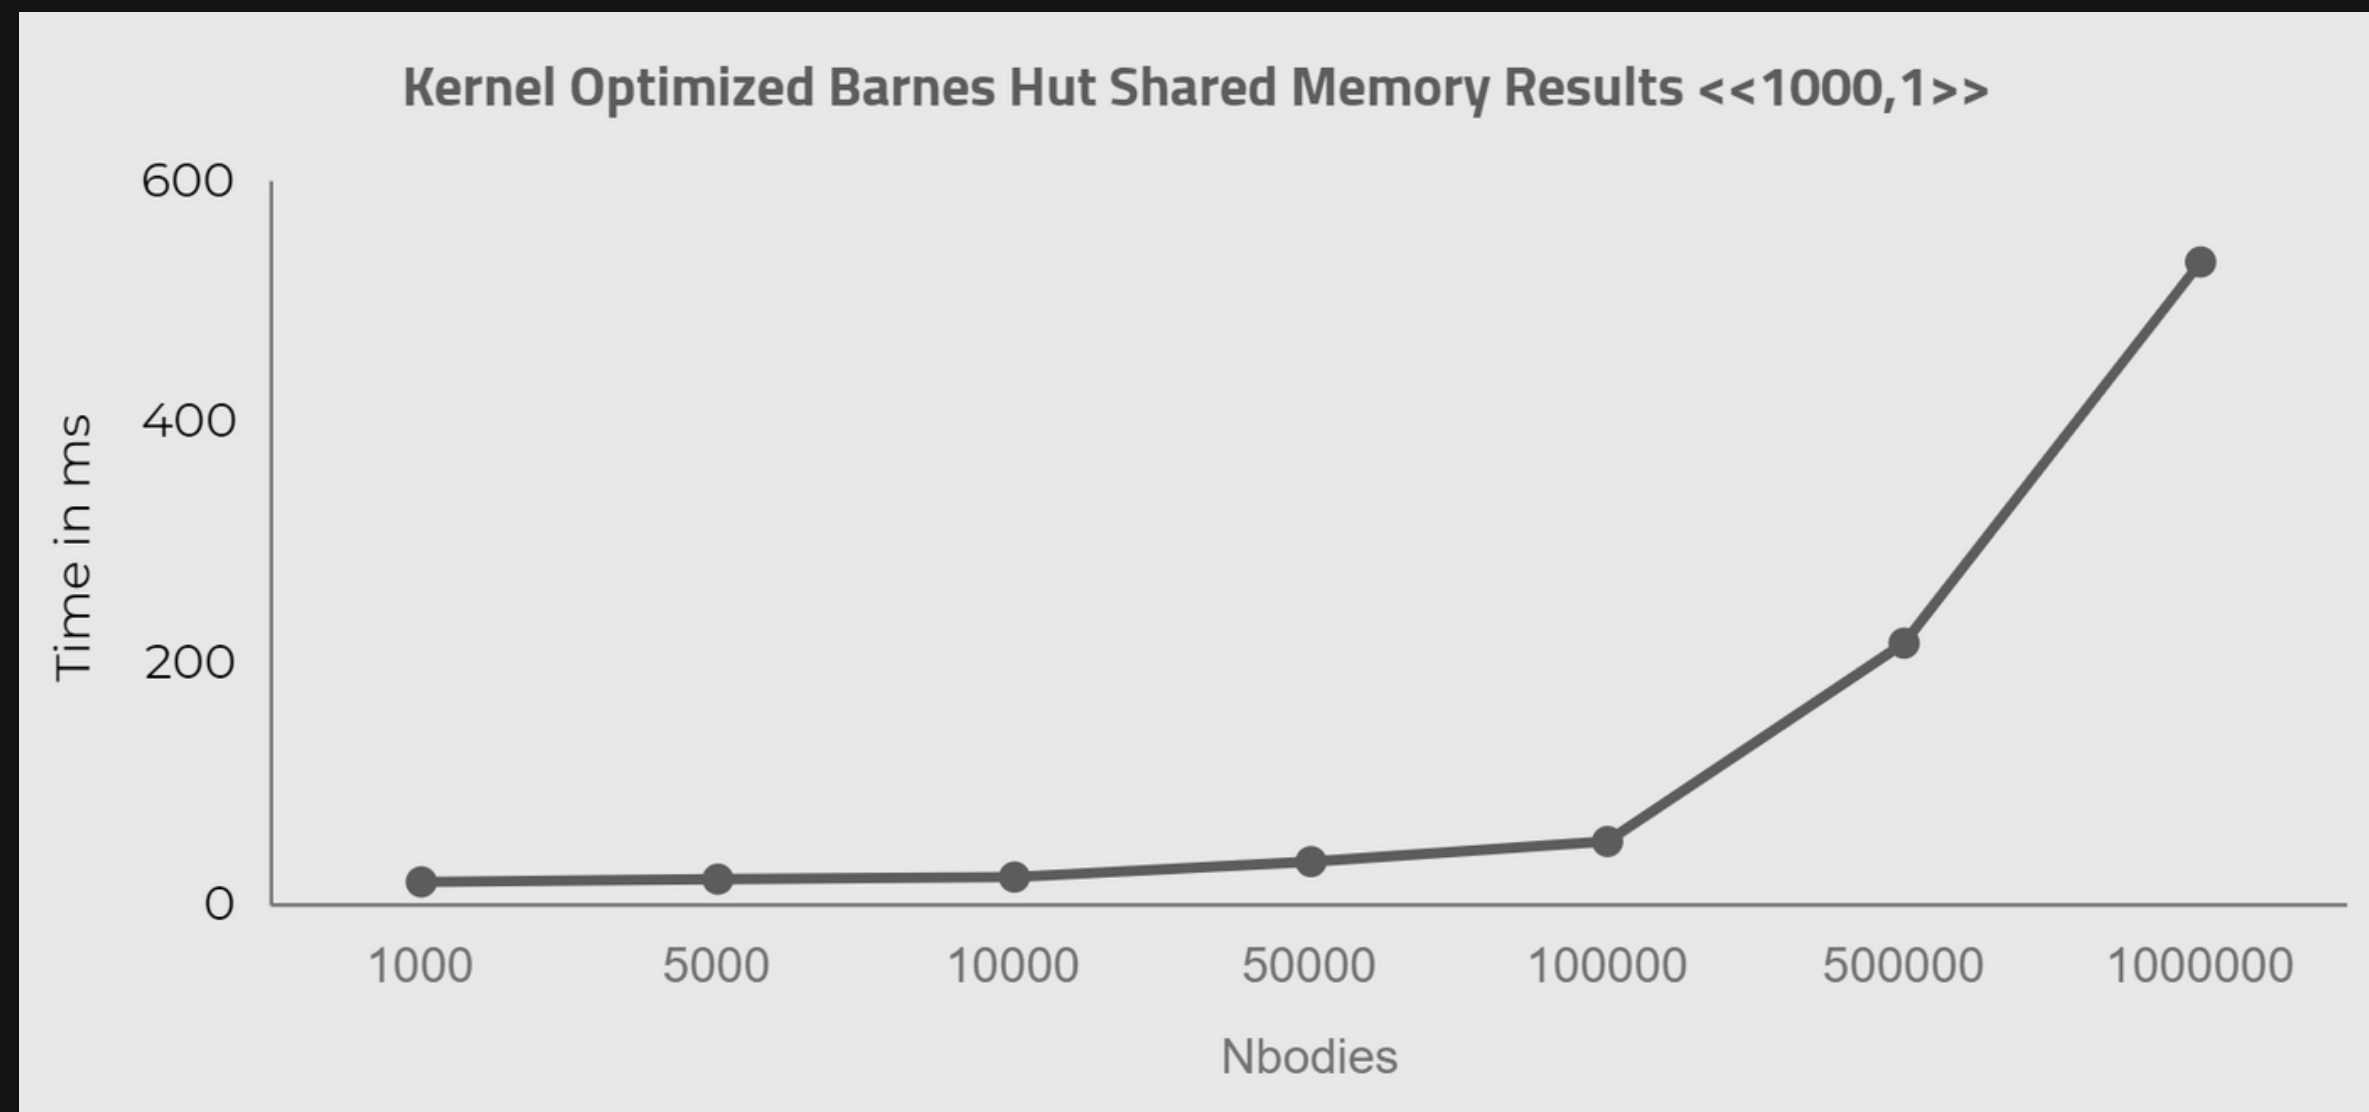


CUDA Shared Memory Implementation Performance




Applied CUDA Optimizations and Performance

INCREASING NUMBER OF BLOCKS FROM 1 TO
1000



Applied CUDA Optimizations and Performance

INCREASING NUMBER OF THREADS FROM 1 TO
1000



Conclusion



Shared Memory



Naive Parallel



Serial

Thank you!!

